

Jolt Formal Specification

Cat

February 02, 2026

1. Expanded Instructions

RISC-V descriptions are taken from <https://msyksphinz-self.github.io/riscv-isadoc/>

1.1. MULH

Done above

1.2. DIV (Special)

Prover: computes q and r (fills advice fields) Prover: constructs a proof that these values satisfy all constraints Verifier: checks the proof

asm

```
1 div rd,rs1,rs2
```

Description perform an XLEN bits by XLEN bits signed integer division of rs1 by rs2, rounding towards zero.

rust

```
1 x[rd] = x[rs1] /s x[rs2]
```

1.2.a. Virutal Instructions

rust

```
1 let a0 = self.operands.rs1; // dividend
2 let a1 = self.operands.rs2; // divisor
3 let a2 = allocator.allocate(); // quotient (from oracle)
4 let a3 = allocator.allocate(); // |remainder| (from oracle)
5 let t0 = allocator.allocate(); // adjusted divisor
6 let t1 = allocator.allocate(); // temporary
7 let shmat = 63
8
9 asm.emit_j::<VirtualAdvice>(*a2, 0); // quotient
10 asm.emit_j::<VirtualAdvice>(*a3, 0); // |remainder|
11
12 // Handle special cases: div-by-zero and overflow
13 // If divisor 0, check q = u64::MAX, otherwise do nothing
14 // rs1 = a1 = divisor
15 // rs2 = a2 = quotient
16 asm.emit_b::<VirtualAssertValidDiv0>(a1, *a2, 0);
17 // t0 = either is divisor (a1) at or 1 (based on overflow)
18 asm.emit_r::<VirtualChangeDivisor>(*t0, a0, a1);
19
20 // Verify no overflow: quotient × divisor must not overflow
```

```

21 // t1 = High64[q * t0]
22 asm.emit_r::<MULH>(*t1, *a2, *t0); // High bits of multiplication
23
24 let t2 = allocator.allocate();
25 let t3 = allocator.allocate();
26
27 // Take lower 64 bits of q*t0 into t2 (always fits in 64 bits)
28 asm.emit_r::<MUL>(*t2, *a2, *t0); // quotient × adjusted_divisor
29 // shmat is 63 - so lower 5 bits are all 1
30 // t3 = sign(t2) = t2>>63
31 asm.emit_i::<SRAI>(*t3, *t2, shmat); // Sign-extend low bits
32 // Check that t1 = t3 or High[q*t0] = t2 >> 63 ???
33 asm.emit_b::<VirtualAssertEQ>(*t1, *t3, 0); // Assert no overflow
34
35 // Apply sign of dividend to remainder
36 asm.emit_i::<SRAI>(*t1, a0, shmat); // Sign bit of dividend
37 asm.emit_r::<XOR>(*t3, *a3, *t1); // XOR with |remainder|
38 asm.emit_r::<SUB>(*t3, *t3, *t1); // Two's complement if negative
39
40 // Verify: dividend = quotient × divisor + remainder
41 asm.emit_r::<ADD>(*t2, *t2, *t3); // Add signed remainder
42 asm.emit_b::<VirtualAssertEQ>(*t2, a0, 0); // Assert equals dividend
43
44 // Verify: |remainder| < |divisor|
45 asm.emit_i::<SRAI>(*t1, *t0, shmat); // Sign bit of adjusted divisor
46 asm.emit_r::<XOR>(*t3, *t0, *t1); // XOR to get magnitude
47 asm.emit_r::<SUB>(*t3, *t3, *t1); // |adjusted_divisor|
48 // checks if (divisor) `t3= 0` or (remainder) `a3 < t3`
49 asm.emit_b::<VirtualAssertValidUnsignedRemainder>(*a3, *t3, 0);
50
51 // Move quotient to destination register
52 asm.emit_i::<ADDI>(<self>.operands.rd, *a2, 0);

```

Claim:

Lemma 1.2.1.1: The assertion $\text{High64}[q \cdot t_0] = (\text{Low64}[q \cdot t_0] \gg 63)$ passes if and only if $q \cdot t_0 \in [-2^{63}, 2^{63} - 1]$.

Proof. \implies :

Suppose $q \cdot t_0 \in [-2^{63}, 2^{63} - 1]$. When viewed as a 128-bit value, this equals the sign extension of its 64-bit representation.

By definition of sign extension:

$$\text{High64} = \begin{cases} 0x0000_0000_0000_0000 & \text{if } q \cdot t_0 \geq 0 \\ 0xFFFF_FFFF_FFFF_FFFF & \text{if } q \cdot t_0 < 0 \end{cases} \quad (1)$$

The arithmetic right shift by 63 positions produces:

$$\text{Low64} \gg 63 = \begin{cases} 0 & \text{if sign bit of Low64 is 0} \\ -1 = 0xFFFF_FFFF_FFFF_FFFF & \text{if sign bit of Low64 is 1} \end{cases} \quad (2)$$

Since the sign bit of Low64 correctly represents whether $q \cdot t_0 \geq 0$ or $q \cdot t_0 < 0$, we have:

$$\text{High64} = \text{Low64} \gg 63 \quad (3)$$

Therefore the assertion passes.

\Leftarrow If the assertion passes, then $q \cdot t_0 \in [-2^{63}, 2^{63} - 1]$

Suppose the assertion $\text{High64} = \text{Low64} \gg 63$ passes.

Case 1: $\text{Low64} \geq 0$ (when interpreted as signed 64-bit, i.e., sign bit = 0)

- Then $\text{Low64} \gg 63 = 0$
- So $\text{High64} = 0 = 0x0000_0000_0000_0000$
- The 128-bit value is: $0x0000_0000_0000_0000 \parallel \text{Low64}$
- This represents a value in $[0, 2^{63} - 1]$ ✓

Case 2: $\text{Low64} < 0$ (when interpreted as signed 64-bit, i.e., sign bit = 1)

- Then $\text{Low64} \gg 63 = -1 = 0xFFFF_FFFF_FFFF_FFFF$
- So $\text{High64} = 0xFFFF_FFFF_FFFF_FFFF$
- The 128-bit value is: $0xFFFF_FFFF_FFFF_FFFF \parallel \text{Low64}$
- This represents Low64 sign-extended to 128 bits
- Since $\text{Low64} \in [-2^{63}, -1]$ as a signed 64-bit value, the 128-bit representation is also in $[-2^{63}, -1]$ ✓

□

Lemma 1.2.1.2: $t_3 = \text{sign}(a)|r|$

Proof. **Step 1:** Extract sign of dividend

$$t_1 = a_0 \gg 63 = \begin{cases} 0 = 0x0000_0000_0000_0000 & \text{if } a \geq 0 \\ -1 = 0xFFFF_FFFF_FFFF_FFFF & \text{if } a < 0 \end{cases} \quad (4)$$

Step 2: XOR with absolute remainder

$$t_3 = a_3 \oplus t_1 \quad (5)$$

Step 3: Subtract to apply two's complement

$$t_3 = t_3 - t_1 \quad (6)$$

Case 1: $a = a_0 \geq 0$ (so $r = a_3$ should be non-negative)

- $t_1 = 0$

- $t_3 = a_3 \oplus 0 = a_3$
- $t_3 = a_3 - 0 = a_3$
- Result: $t_3 = |r| = r$

Case 2: $a < 0$ (so $r = a_3$ should be ≤ 0)

- $t_1 = -1 = 0xFFFF_FFFF_FFFF_FFFF$
- $t_3 = a_3 \oplus 0xFFFF_FFFF_FFFF_FFFF = \tilde{a}_3$ (bitwise NOT)
- $t_3 = \tilde{a}_3 - (-1) = \tilde{a}_3 + 1$
- By two's complement: $\tilde{a}_3 + 1 = -a_3 = -|r|$
- Result: $t_3 = -|r|$

□

1.2.b. The Code (last two conditons are quite easy)

rust

```
1 // Verify: dividend = quotient × divisor + remainder
2 asm.emit_r::<ADD>(*t2, *t2, *t3);      // t2 = q·b + r
3 asm.emit_b::<VirtualAssertEQ>(*t2, a0, 0); // Assert equals dividend
```

1.2.c. What This Checks

At this point in the code:

- $t_2 = q \cdot t_0$ where t_0 is the adjusted divisor
- $t_3 = r$ with correct sign
- $a_0 = a$ (dividend)

The assertion verifies: $a = q \cdot b + r$

1.2.d. Proof of Correctness

Normal case: $t_0 = b$ (no overflow)

- The code computes: $t_2 + t_3 = q \cdot b + r$
- The assertion checks: $q \cdot b + r = a$
- This is exactly the division algorithm requirement ✓

Overflow case: $(a, b) = (-2^{63}, -1)$, and $r = 0$ from the next condition that $|r| < |b'|$ (Adjusted)

- To avoid computing $q \cdot b = (-2^{63}) \cdot (-1) = 2^{63}$ (overflow)
- The code sets $t_0 = 1$ (adjusted divisor)
- Then the check becomes $q + r = (-2^{63}) \cdot 1 = -2^{63}$
- If the prover was honest $r = 0$ and $q = -2^{63}$ is the correct answer, and passes. The assertion verifies: $q \cdot t_0 + r = -2^{63} + 0 = a$ ✓

Theorem 1.2.4.1: State matches.

Proof. All we need to show is that the final value in `rd` which we denote with z is actually $\frac{x}{y}$ signed integer division with rounding towards 0. We start with first two

instructions which simply stores in virtual registers `a2` and `a3` claimed quotient q , and remainder r .

Next, we focus on the last instruction which simply returns q as the answer. The thing to show here, if the prover uses a bad value for q and r , one that is not what the real RISC-V cpu would have computed, we would have a panic, and thus the program would crash. \square

⚠ Warning

Note for ARI:

The guest program says `div`; and then i expand that to above. Now notice that the program starts with two instrucions.

```
rust
1 asm.emit_j::<VirtualAdvice>(*a2, 0); // quotient
2 asm.emit_j::<VirtualAdvice>(*a3, 0); // |remainder|
```

These two instructions are telling the CPU, write into virtual registers `a2` and `a3` the quotient and remainder. Now where does the CPU, get these answers from? Well it gets it from somewhere, as q and r which it then proceeds to put into these registers. The code below shows what the CPU actually does during tracing.

```
rust
1 let mut inline_sequence = self.inline_sequence(&cpu.vr_allocator,
    cpu.xlen);
2 // The first instructioon
3 if let Instruction::VirtualAdvice(instr) = &mut inline_sequence[0] {
4     instr.advice = quotient;
5 } else {
6     panic!("Expected Advice instruction");
7 }
8 // The Second instruction
9 if let Instruction::VirtualAdvice(instr) = &mut inline_sequence[1] {
10     instr.advice = remainder;
11 } else {
12     panic!("Expected Advice instruction");
13 }
14
15 // With these values in there it just executes as normal.
16 let mut trace = trace;
17 for instr in inline_sequence {
18     instr.trace(cpu, trace.as_deref_mut());
19 }
```

1.3. SUBW

```
asm
1 subw rd,rs1,rs2
```

Description: Subtract the 32-bit of registers rs1 and 32-bit of register rs2 and stores the result in rd. Arithmetic overflow is ignored and the low 32-bits of the result is sign-extended to 64-bits and written to the destination register

1.3.a. Virtual Sequence

rust

```
1 asm.emit_r::<SUB>(&self.operands.rd, &self.operands.rs1,
                      &self.operands.rs2);
2 asm.emit_i::<VirtualSignExtendWord>(&self.operands.rd, &self.operands.rd,
                                         0);
```

where the instructions SUB and VirtualSignExtendWord are defined as follows:

SUB: Subs the register rs2 from rs1 and stores the result in rd. Arithmetic overflow is ignored and the result is simply the low XLEN bits of the result.

rust

```
1 //Implementation
2 x[rd] = x[rs1] - x[rs2]
```

VirtualSignExtendWord: Sign-extends the lower 32 bits of a register to 64 bits:

rust

```
1 //Implementation
2 (cpu.x[&self.operands.rs1 as usize] << 32) >> 32
```

Theorem 1.3.1.1: Let the triple $(rd, rs1, rs2)$ denote the machine state. Then machine state before and after execution of subw and the virtual sequence is identical.

Proof. $rs1, rs2$ remain unchanged before and after the instruction. All there is to show is that the value in rd is the same.

As defined earlier let x, y, z denote the w bit values in $rs1, rs2, rd$ represented in unsigned form. Similarly, x', y', z' are the same but in signed representation.

The first instruction gives us :

$$z = x - y \pmod{2^{64}} \quad (7)$$

By construction, the lower 32 bits of z are exactly with arithmetic overflow ignored.

$$z[31 : 0] = (x[31 : 0] - y[31 : 0]) \pmod{2^{32}} \quad (8)$$

Finally we sign extend the lower 32 bits of z to get the desired answer. \square

1.4. SLLI

asm

```
1 slli rd,rs1,shamt
```

Performs logical left shift on the value in register `rs1` by the shift amount held in the lower 5 bits of the immediate. In RV64, bit-25 is used to `shamt[5]`.

1.4.a. Virtual Sequence

rust

```
1 shift = imm & 0x3f; // Lower 6 bits
2 asm.emit_i::<VirtualMULI>(&self.operands.rd, &self.operands.rs1, 1 <<
    shift);
```

Theorem 1.4.1.1: States mactch

Proof.

Let variable s denote the `shift` amount stored in `Imm[0:5]` and x denote the value in `rs1`. It's a mathematical fact that

$$z = x \times 2^s = x \ll s \quad (9)$$

where \times denotes wrapping multiplication i.e we drop the overflow. Therefore, the final value in `rd` is correct. \square

1.5. SRLI

asm

```
1 rd, rs1, imm
```

Description: Performs logical right shift on the value in register `rs1` by the shift amount held in the lower 5 bits of the immediate In RV64, bit-25 is used to `shamt[5]`.

1.5.a. Virtual Sequence

rust

```
1 shift = imm & 0x3f // shift \in [0, 63]
2 len = 64
3 ones = (1u128 << (len - shift)) - 1;
4 bitmask = (ones << shift) as u64
5 // Virtual SRLI right shifts contents of rs1 by the number
6 // of trailing zeros in the bitmask
7 asm.emit_vshift_i::<VirtualSRLI>(&self.operands.rd, &self.operands.rs1,
    bitmask);
```

Theorem 1.5.1.1: Match

Proof. Let i denote the immediate value in the instruction. Define $w = 64$ and $s = i \wedge 0x3F$ be the number of bits to shift.

From the definition of `srli` we have that the destination register will contain

$$z = x \gg s \quad (10)$$

Define $o := 2^{w-s} - 1$. We claim o has s 0's in its higher order bits. If this claim were to be true, then the bitmask b created by logically shifting o by s is given by $b := o \ll s$, and would have exactly s trailing 0's. We know that `VirtualSRLI` right shifts contents of `rs1` by number of trailing 0's of `bitmask`. This gives us that $z = x \gg s$.

All that's left to show is that b has s 0's in its higher order bits. For any $k \geq 0$, 2^k has bit k set to 1, and k trailing 0's, and $w-k-1$ prefix bits are 0. This gives us $2^k - 1$, clears the k 'th bit, and sets all trailing 0's to ones, giving us $w-k$ prefix bit 0's. Now set $k = w-s$, we get our result.

□

1.6. LW

asm

```
1 lw rd,offset(rs1)
```

Description: Loads a 32-bit value from memory and sign-extends this to XLEN bits before storing it in register rd.

rust

```
1 x[rd] = sext(M[x[rs1] + sext(offset)][31:0])
```

1.6.a. Virtual Sequence

rust

```
1 // Check if x[rs1] + imm % 4 == 0 panic otherwise
2 asm.emit_halign::<VirtualAssertWordAlignment>(&self.operands.rs1,
    &self.operands.imm);
3 // v_address = x[rs1] + imm[0:11] with wrapping add (this is the address
   from which i want to load)
4 asm.emit_i::<ADDI>(*v_address, &self.operands.rs1, &self.operands.imm as
   u64);
5 // v_dword_address = v_address & (-8 as i64) as u64
6 asm.emit_i::<ANDI>(*v_dword_address, *v_address, -8i64 as u64);
7 // v_dword = M[v_dword_address]; which is v_address with lowest 3 bits
   cleared
8 asm.emit_ld::<LD>(*v_dword, *v_dword_address, 0);
9 // v_shift = v_address << 3
10 asm.emit_i::<SLLI>(*v_shift, *v_address, 3);
11 // rd = v_d >> v_shift[0:4] ; r-shift x[rs1] amount held in the lower 5
   bits of v_shift
12 asm.emit_r::<SRL>(&self.operands.rd, *v_dword, *v_shift);
13 // rd = sign-extend(rd)
14 asm.emit_i::<VirtualSignExtendWord>(&self.operands.rd, &self.operands.rd,
   0);
```

Note for Ari: this proof can be shortened, but i overgenerated so i remember.

Proof.

css

1 Address (hex)	LB	LH	LW	LD	
2 0x1000	✓	✓	✓	✓	<- this is now v_(dwa)
3 0x1001	✓	✗	✗	✗	
4 0x1002	✓	✓	✗	✗	
5 0x1004	✓	✓	✓	✗	<- pretend this was v_a = v_(wa)
6 0x1008	✓	✓	✓	✓	

As we will load 32 bits or 1 word into memory, we need our address to be word aligned i.e divisible by 4. This is the first check.

1. We first compute the address $a := x + i$, where $i = \text{imm}$, and check that $a \bmod 4 = 0$.
2. If the above test passes set $v_{wa} := a$ where a is defined above.
3. Clear the lowest 3 bits of v_{wa} and set $v_{dwa} = v_{(wa)'} \parallel 000$ where $v_{(wa)'}$ is the highest $w - 3$ bits of v_{wa} .

We do this because we need our temporary double word address to be divisible by 8, or double-word aligned.

4. Our double word aligned address is good for a load; $v_d := M[v_{dwa}]$; load 64 bits (2 words) from address in register v_{dwa}
5. $v_s := v_a \ll 3$

This is subtle. Note that we are guaranteed that v_a is word-aligned, but it could also be double-word aligned. In v_d we have the 64 bits read from location v_a . If v_a were also double word aligned, then it would mean that $v_a = v_{wa} = v_{dwa}$. So we want the lower 32 bits of v_d as the answer. Instead, if we v_a was not double word aligned, then we want the high 32 bits of v_d as the answer (as shown in the example above).

The way SRL works is it will look at the lower 5 bits of v_s to decide how much to shift right by. If we are in the first case, and v_a is double word aligned, then it already has 3 trailing 0's. Now v_s will have 6 trailing 0's, and we are safe that we do not right shift v_d at all. We just output the lower 32 bits sign extended as the answer. If v_a were not double word aligned, then we are guaranteed that there should be a 1 in index 2. Now if we were to left shift 3 units, we'd have $v_s = 32$, and this gives us the higher 32 bits. This completes the proof.

□

1.7. SRA

asm

```
1 sra rd,rs1,rs2
```

Description Performs arithmetic right shift on the value in register rs1 by the shift amount held in the lower 5 bits of register rs2

Implementaion:

rust

```
1 x[rd] = x[rs1] >> s x[rs2]
```

1.7.a. Virtualisation

rust

```
1 // v_bitmask = bitmask such that it has s=x[rs2][0:5] trailing 0s
2 asm.emit_i::<VirtualShiftRightBitmask>(*v_bitmask, self.operands.rs2, 0);
3 // Right shift x[rs1] by num trailing 0's in v_bitmask and store in rd
4 asm.emit_vshift_r::<VirtualSRA>(self.operands.rd, self.operands.rs1,
    *v_bitmask);
```

Proof.

1. The first instruction constructs v_b with s trailing 0's, where $s = y[0 : 5]$.
2. Now we right shift x by s

The equivalence holds trivially. \square

1.8. AMONEXW

asm

```
1 amomaxu.w rd, rs2, (rs1)
```

Description Atomically load a 32-bit unsigned data value from the address in rs1, place the value into register rd, apply unsigned max the loaded value and the original 32-bit unsigned value in rs2, then store the result back to the address in rs1.

rust

```
1 x[rd] = AM032(M[x[rs1]] MAXU x[rs2])
```

1.8.a. Virtual Instructions

rust

```
1
2 asm.emit_halign::<VirtualAssertWordAlignment>(rs1, 0);
3 // Use v_shift temporarily to hold aligned address
4 asm.emit_i::<ANDI>(v_shift, rs1, -8i64 as u64);
5 asm.emit_ld::<LD>(v_dword, v_shift, 0);
6 // Now compute the actual shift value
7 asm.emit_i::<SLLI>(v_shift, rs1, 3);
8 asm.emit_r::<SRL>(v_rd, v_dword, v_shift);
9
10
11 asm.emit_i::<VirtualZeroExtendWord>(*v_rs2, self.operands.rs2, 0);
12 // Zero-extend v_rd in place into v_sel_rd (temporarily)
13 asm.emit_i::<VirtualZeroExtendWord>(*v_sel_rd, *v_rd, 0);
14 // Compare: v_sel_rd (zero-extended v_rd) < v_rs2
15 asm.emit_r::<SLTU>(*v_sel_rs2, *v_sel_rd, *v_rs2);
16 // Invert selector to get selector for v_rd
17 asm.emit_i::<XORI>(*v_sel_rd, *v_sel_rs2, 1);
18 // Select maximum using multiplication
19 asm.emit_r::<MUL>(*v_rs2, *v_sel_rs2, self.operands.rs2);
20 asm.emit_r::<MUL>(*v_sel_rd, *v_sel_rs2, *v_rd);
```

```

21 asm.emit_r::<ADD>(*v_rs2, *v_sel_rd, *v_rs2);
22
23
24 asm.emit_i::<ORI>(v_mask, 0, -1i64 as u64);
25 asm.emit_i::<SRLI>(v_mask, v_mask, 32);
26 asm.emit_r::<SLL>(v_mask, v_mask, v_shift);
27 // Use v_shift as temporary after it's been used for shifting
28 asm.emit_r::<SLL>(v_shift, rs2, v_shift);
29 asm.emit_r::<XOR>(v_shift, v_dword, v_shift);
30 asm.emit_r::<AND>(v_shift, v_shift, v_mask);
31 asm.emit_r::<XOR>(v_dword, v_dword, v_shift);
32 // Recompute aligned address for store
33 asm.emit_i::<ANDI>(v_mask, rs1, -8i64 as u64);
34 asm.emit_s::<SD>(v_mask, v_dword, 0);
35 asm.emit_i::<VirtualSignExtendWord>(rd, v_rd, 0);

```

1.9. SRLIW

1 `slliw rd,rs1,shamt`

asm

Description Performs logical left shift on the 32-bit of value in register rs1 by the shift amount held in the lower 5 bits of the immediate. Encodings with imm[5] ≠ 0 are reserved.

The virtual instructions are as follows:

1 `asm.emit_i::<SLLI>(*v_rs1, self.operands.rs1, 32);`
2 `(shift, len) = ((self.operands.imm & 0x1f) + 32, 64),`
3 `let ones = (1u128 << (len - shift)) - 1;`
4 `let bitmask = (ones << shift) as u64;`
5 `asm.emit_vshift_i::<VirtualSRLI>(self.operands.rd, *v_rs1, bitmask);`
6 `asm.emit_i::<VirtualSignExtendWord>(self.operands.rd, self.operands.rd, 0);`

rust

Theorem 1.9.1: Match

Proof. Want to show that the above virtual instructions performs $rd = (rs1 \text{ as } u32) \gg shamt$ where $s=shamt = \text{imm}[4:0]$

Let x, z denote the values in $rs1, rd$ respectively. Let v_1 denote the contents of the virtual register v_{rs1} . Based in the first instruction we have v_1 holds the lower 32 bits of x in its upper 32 bits.

$$v_1 = x[0 : 31] \parallel 0 \quad (11)$$

Define $s' := s + 32$ where s is the specified shift amount. This implies $s' \in [32, 63]$, as $s \in [0, 31]$.

Define

$$o := 2^{64-(s+32)} - 1 = 2^{32-s} - 1 \quad (12)$$

Here o is a w bit integer with exactly $(32 - s)$ lower bits set to 1. Therefore the $32 + s$ upper bits are set to 0.

$$b := (o \ll s') = (o \ll (s + 32)) \quad (13)$$

Now b has $s + 32$ trailing 0's so when we right shift v_1 by $s + 32$ bits, the 32 first ensures that $v_1 := 0 \parallel x[0 : 31]$, so it's x as u32, and then it right shifts by s bits as desired. \square

1.10. SRL

asm

```
1 srl rd, rs1, rs2
```

Description: Logical right shift the contents of **rs1** by **shamt**, where **shamt** lower 6 bits of **rs2** as a unsigned integer.

1.10.a. Virtual Sequence

rust

```
1 asm.emit_i::<VirtualShiftRightBitmask>(*v_bitmask, self.operands.rs2, 0);
2 asm.emit_vshift_r::<VirtualSRL>(self.operands.rd, self.operands.rs1,
 *v_bitmask);
```

Theorem 1.10.1.1: Match

Proof.

1. Let s denote the lower 6 bits of **rs2**. The first instruction sets **v_bitmask** to w bit number with s trailing 0's, from the guarantees of **VirtualShiftRightBitmask**.
2. The second instruction logically right shifts the contents of **rs1** by s (the number of trailing 0's in **v_bitmask**).

\square

1.11. SW

Store at memory location in **rs1** the lower 32 bits of the contents of **rs2**.

rust

```
1
2 asm.emit_halign::<VirtualAssertWordAlignment>(self.operands.rs1,
 self.operands.imm);
3 asm.emit_i::<ADDI>(*v_address, self.operands.rs1, self.operands.imm as
 u64);
4 asm.emit_i::<ANDI>(*v_dword_address, *v_address, -8i64 as u64);
5 asm.emit_ld::<LD>(*v_dword, *v_dword_address, 0);
6 asm.emit_i::<SLLI>(*v_shift, *v_address, 3);
```

```

7 asm.emit_i::<ORI>(*v_mask, 0, -1i64 as u64);
8 asm.emit_i::<SRLI>(*v_mask, *v_mask, 32);
9 asm.emit_r::<SLL>(*v_mask, *v_mask, *v_shift);
10 asm.emit_r::<SLL>(*v_word, self.operands.rs2, *v_shift);
11 asm.emit_r::<XOR>(*v_word, *v_dword, *v_word);
12 asm.emit_r::<AND>(*v_word, *v_word, *v_mask);
13 asm.emit_r::<XOR>(*v_dword, *v_dword, *v_word);
14 asm.emit_s::<SD>(*v_dword_address, *v_dword, 0);

```

Theorem 1.11.1: Match

Proof.

We first check if `v_address` is word aligned or not. If not we panic immediately, as the original specification would. If it were to be aligned we have two cases:

1. `v_address` is also double word aligned i.e divisible by 8.
2. `v_address` is only single word aligned but not double word aligned i.e. not divisible by 8.

`asm.emit_i::<ANDI>(*v_dword_address, *v_address, -8i64 as u64);`
`v_dword_address` is by definition double word aligned as we AND it with a mask with 3 trailing 0s.

`asm.emit_ld::<LD>(*v_dword, *v_dword_address, 0);` `v_dword` stores 64 bits starting at address `v_dword_address`

rust

```
1 asm.emit_i::<SLLI>(*v_shift, *v_address, 3);
```

Based on whether we are in case 1 or case 2, `v_shift=0` or `v_shift=32`

rust

```
1 asm.emit_i::<ORI>(*v_mask, 0, -1i64 as u64);
```

This sets `v_mask` to the all ones bits string.

`asm.emit_i::<SRLI>(*v_mask, *v_mask, 32);` This clears the upper 32 bits of the mask.

`asm.emit_r::<SLL>(*v_mask, *v_mask, *v_shift);` Now if we are in case 1, `v_mask=0 || 1` were upper 32 bits are 0, and lower 32 bits are 1. If we are in case 2, `v_mask=1 || 0` were upper 32 bits are 1, and lower 32 bits are 0.

`asm.emit_r::<SLL>(*v_word, self.operands.rs2, *v_shift);` If we are in case 1, then `v_word = x[rs2]` If we are in case 2, then `v_word = x[rs2][0:31] || 0`, i.e the lower 32 bits of `rs2` are the upper 32 bits of `v_word`.

Let `mem = v_dword`.

```
asm.emit_r::<XOR>(*v_word, *v_dword, *v_word); In case 1, v_word = x[rs2] XOR  
mem In case 2, v_word = mem[63:32] || (mem[0:31] XOR x[rs2][0:31])
```

```
asm.emit_r::<AND>(*v_word, *v_word, *v_mask); If case 1, where v_word was the  
entire contents of rs2, and v_mask= 0 ||1, this will set v_word to just 0 || x[rs2]  
[0:31] XOR mem[0:31].
```

If case 2, where $v_Word = x[rs2][0:31] \text{ XOR } \text{mem}[63:32] \mid\mid 0$, and $v_mask = 1 \mid\mid 0$, this will keep v_word as is.

```
asm.emit_r::<XOR>(*v_dword, *v_dword, *v_word);
```

```
Now in case 1, v_dword = mem[63:32] || x[rs2][31:0] XOR mem[0:31] XOR mem[31:0]  
= mem[63:32] || x[rs2][31:0], exactly what I want. By similar, analysis, In case 2,  
v_dword = x[rs2][31:0] || mem[31:0].
```

Store it in the right location, and that completes the proof.

```
asm.emit_s::<SD>(*v_dword_address, *v_dword, 0);
```

□

1.12. REMU

If the quotient and remainder are correct, then we want to show all the checks will pass and the final value in $rd = a3$.

Let r be the remainder the native RISC-V instruction REMU returns. We are guaranteed that there exists some q such that $a0 = a1 \times q + r$ where $0 \leq r < a1$ (folklore number theory fact). Assume that the prover sets $a2=q$ and $a3=r$.

Then, $a2 \times a1$ should not overflow, as $a0 = a1 \times a2 + a3$ and $a0 \in [0, 2^{64} - 1]$, then both the following check passes.

rust

```
1 asm.emit_b::<VirtualAssertMulUNoOverflow>(*a2, a1, 0)  
2 asm.emit_b::<VirtualAssertEQ>(*t1, a0, 0)
```

`asm.emit_b::<VirtualAssertLTE>(*t0, *t1, 0)` checks $a1 \times a2 \leq a1 \times a2 + a3$ which passes as $r=a3 \geq 0$,

`asm.emit_b::<VirtualAssertLTE>(*a3, *t1, 0)` checks if $a3 \leq a1 \times a2 + a3$ which also holds with equality. All tests pass and rd is correct. The contrapositive, says that if one of the tests fails, then either $a2$ and $a3$ are different from q and r .

1.13. AMOSWAPW

asm

```
1 TODO
```

Description: Load a word from `address`, where `address = x[rs1] + offset`, sign extend it and put it in `rd`. Also put the lower 32 bits of the contents of `rs2` in at the memory location `address`.

1.13.a. Virtual Sequence

rust

```
1 asm.emit_halign::<VirtualAssertWordAlignment>(&self.operands.rs1, 0);
2 asm.emit_i::<ANDI>(*v_shift, &self.operands.rs1, -8i64 as u64);
3 asm.emit_ld::<LD>(*v_dword, *v_shift, 0);
4 asm.emit_i::<SLLI>(*v_shift, &self.operands.rs1, 3);
5 asm.emit_r::<SRL>(*v_rd, *v_dword, *v_shift);
6 asm.emit_i::<ORI>(*v_mask, 0, -1i64 as u64);
7 asm.emit_i::<SRLI>(*v_mask, *v_mask, 32);
8 asm.emit_r::<SLL>(*v_mask, *v_dword, *v_shift);
9 asm.emit_r::<SLL>(*v_shift, &self.operands.rs2, *v_shift);
10 asm.emit_r::<XOR>(*v_shift, *v_dword, *v_shift);
11 asm.emit_r::<AND>(*v_shift, *v_dword, *v_mask);
12 asm.emit_r::<XOR>(*v_dword, *v_dword, *v_shift);
13 asm.emit_i::<ANDI>(*v_mask, &self.operands.rs1, -8i64 as u64);
14 asm.emit_s::<SD>(*v_mask, *v_dword, 0);
15 asm.emit_i::<VirtualSignExtendWord>(&self.operands.rd, *v_rd, 0);
```

1.13.b. Proof Of Correctness

Proof on whiteboard.

1.14. amoord

Description: Atomically loads a 64-bit doubleword from the address in `rs1`, ORs it with `rs2`, stores the result back at the same address. Place the original value loaded from the address value in `rd`.

1.14.a. Virtual Sequence

rust

```
1 asm.emit_ld::<LD>(*v_rd, &self.operands.rs1, 0);
2 asm.emit_r::<OR>(*v_rs2, *v_rd, &self.operands.rs2);
3 asm.emit_s::<SD>(&self.operands.rs1, *v_rs2, 0);
4 asm.emit_i::<ADDI>(&self.operands.rd, *v_rd, 0);
```

1.14.b. Proof Of Correctness

Lines 1-2 set `v_rs2 = (M[x[rs1]] OR x[rs2])` and Line 3 stores `v_rs2` at memory location stored in `rs1`. Line 4 puts `M[x[rs1]]` into `rd`.

1.15. LBU

Description: Loads an 8-bit byte from memory at `rs1 + offset`, zero-extending the result into `rd`.

1.15.a. Virtual Sequence

rust

```
1 asm.emit_i::<ADDI>(*v_address, &self.operands.rs1, self.operands.imm as
u64);
2 asm.emit_i::<ANDI>(*v_dword_address, *v_address, -8i64 as u64);
3 asm.emit_ld::<LD>(*v_dword, *v_dword_address, 0);
```

```

4 asm.emit_i::<XORI>(*v_shift, *v_address, 7);
5 asm.emit_i::<SLLI>(*v_shift, *v_shift, 3);
6 asm.emit_r::<SLL>(self.operands.rd, *v_dword, *v_shift);
7 asm.emit_i::<SRLI>(self.operands.rd, self.operands.rd, 56);

```

1.15.b. Proof Of Correctness

1.16. SLLW

Description: Shifts the lower 32 bits of rs1 left by the amount in lower 5 bits of rs2, sign-extending the 32-bit result into rd.

1.16.a. Virtual Sequence

rust

```

1
2 asm.emit_i::<VirtualPow2W>(*v_pow2, self.operands.rs2, 0);
3 asm.emit_r::<MUL>(self.operands.rd, self.operands.rs1, *v_pow2);
4 asm.emit_i::<VirtualSignExtendWord>(self.operands.rd, self.operands.rd,
0);

```

1.16.b. Proof Of Correctness

Left shift is by definition multiplication by a power of 2. The lowest 5 bits of rs2 should be a number between 0 and 31.

By the guarantees of `VirtualPow2W`,

rust

```
1 v_pow2 = 1 << (x[rs2] % 32) = 2^{x[rs2]}
```

Line 2, ensures $rd = v_pow2 \times x[rs1]$, where `mul` Multiplies `rs1` by `rs2` and stores the lower `XLEN` bits of the product in `rd`. Then Line 3, sign-extends the lower 32 bits to 64 bits by the guarantees of `VirtualSignExtendWord`.

1.17. LH

Description: Loads a 16-bit halfword from memory at `rs1 + offset`, sign-extending the result into `rd`.

1.17.a. Virtual Sequence

rust

```

1 asm.emit_halign::<
2 VirtualAssertHalfwordAlignment,
3   >(self.operands.rs1, self.operands.imm);
4 asm.emit_i::<ADDI>(*v_address, self.operands.rs1, self.operands.imm as
u64);
5 asm.emit_i::<ANDI>(*v_dword_address, *v_address, -8i64 as u64);
6 asm.emit_ld::<LD>(*v_dword, *v_dword_address, 0);
7 asm.emit_i::<XORI>(*v_shift, *v_address, 6);
8 asm.emit_i::<SLLI>(*v_shift, *v_shift, 3);
9 asm.emit_r::<SLL>(self.operands.rd, *v_dword, *v_shift);
10 asm.emit_i::<SRAI>(self.operands.rd, self.operands.rd, 48);

```

1.17.b. Proof Of Correctness